

AWI-ESM3 ↔ PISM moving-cavity coupling

the investigation log

Single living document. Superseded claims are struck through, not deleted.

Last updated: 2026-07-13

Abstract

Where we are (2026-07-14, after the fixes). Two real, independent bugs explain the catastrophes of the past week; both are now fixed and under test (movcav16):

Bug I — PISM’s ocean forcing was surface water. The `-ocean` route averaged FESOM level indices 1–20 = **0–410 m including the summer surface**: PISM received ice-base water of up to +6.5°C even from a healthy ocean, melted accordingly (~30 m/yr), and shed **24% of its floating area per decade** — so every mesh-change leg had to absorb a catastrophic geometry increment. FESOM’s *own* cavity melt in the same healthy runs is sane: **2.2 m/yr mean, 0.72 median**. Fix: PISM now consumes FESOM’s own melt directly (`-ocean given`, the DIRECT route; esm_tools 93a207ee).

Bug II — the offline feom exchange weights were in the wrong node ordering. FESOM exchanges coupling fields in *partition (my_list/rank) order*; the offline weight engine (2026-06-17) built them in *nod2d.out order*. Result: every A096↔feom exchange in every submesh leg was **geographically tile-translocated** — the atmosphere was coupled to a jigsaw ocean (Fig. 1): OIFS saw Weddell ice fraction **0.04** while FESOM had **0.95**; global sea-ice death followed within months. Fix: permute to runtime ordering + a mandatory weight validator that aborts the leg on ordering mismatch (c21fe7c2).

Bug III — the OASIS coupling restart was compound-scrambled (found 2026-07-14 by the movcav16 attempts themselves, §6): `remap_oasis_restart.py` gathered in *nod2d* space but its input and required output are *runtime*-ordered, so OIFS’s first coupling interval saw a jigsaw SST — a one-hour global shock that nondeterministically detonated the convection scheme (overflow in `cubasen.F90` at hour 19 / hour 458 / never, on identical inputs). Fix: un-permute / re-permute around the gather via both meshes’ `rpart.out` (abe56ee2); pool `rstos.nc` rebuilt (measured 0.15 → 0.995 correlation with the restart’s own surface temperature).

Four earlier root-cause claims are falsified and struck through below. Sections whose measurements were made on jigsaw-coupled runs are marked **TAINTED** and will be re-measured under movcav16.

Contents

1	The two real bugs (2026-07-14)	2
1.1	Bug I: PISM was fed surface water as ice-base forcing	2
1.2	Bug II: the exchange weights scrambled the coupled ocean	2
1.3	How the two bugs interlocked	4
2	Blocker I — OIFS step-0 coastline crash (RESOLVED)	4
3	Blocker II — FESOM cavity-margin blowup	4
3.1	Two root causes we got wrong	4
3.2	Every hypothesis, and the experiment that killed it	4

3.3	The decisive control: a cold start on PISM’s cavity is stable	5
3.4	What is actually wrong: the dynamics seam	5
4	The increment test (2026-07-13) — TAINTED, to be re-run	5
4.1	The finding: PISM is collapsing	6
4.2	This is not good news	7
5	Falsified: the SPP claim	7
6	Where to go from here	7
6.1	movcav16 launch log (2026-07-14) — and Bug III	8
7	Bugs found and fixed along the way	9
7.1	In the remap	9
7.2	In the coupling plumbing	9
8	Assets	10
9	Method note	10

1 The two real bugs (2026-07-14)

1.1 Bug I: PISM was fed surface water as ice-base forcing

`coupling_ocean2pism.functions` built PISM’s `-ocean th` forcing as a vertical mean over FESOM level *indices* 1–20 (`-sellevidx,1/20`) = **0–410 m, surface included**; the correct `-sellevel,150/600` sat in dead code behind `if [ 0 -eq 1 ]`. Measured consequence in the *healthy* chain (movcav8): `theta_ocean` fed to PISM had mean -0.43°C and **max** $+6.53^{\circ}\text{C}$ (realistic ice-base: ≈ -1.9). PISM’s three-equation scheme melted from that at $\sim 30\text{ m/yr}$ and collapsed the shelves: -24% floating area per decade, 1399 nodes losing their ice *entirely* per 10-yr leg, Ross front -440 km (Fig. ??). Meanwhile **FESOM’s own melt** in those same healthy runs: **2.2 m/yr mean, 0.72 median, deep-cavity 2.6** — sane.

Fix (93a207ee): the DIRECT route. PISM consumes FESOM’s own cavity fluxes via `-ocean given: fw→shelfbmassflux` ($\times 1000$; sign preserved — FESOM melt is $fw < 0$ and PISM’s `shelfbmassflux` is “positive if ice is freezing on”), `sst(level 1 = ice base under the virtual-column convention)→shelfbtemp`. One model owns the ice–ocean interface. The `-ocean given` machinery existed all along; it silently fell back to `th` because its input is a FESOM1.x-era file (`EXP.year.pism.nc`) that FESOM2 never writes — `fesom2ice` now synthesizes it from the standard monthly streams. The 150–600 m band is also enabled as a safety for anyone still on `th`.

1.2 Bug II: the exchange weights scrambled the coupled ocean

FESOM registers its OASIS partition rank-contiguously (Apple partition, `cpl_driver.F90`): the exchange field is the concatenation of each rank’s owned nodes in `my_list` order — **not** `nod2d.out` order (median displacement between the two: 58° of latitude). The offline weight engine introduced on 2026-06-17 built the feom geometry in `nod2d` order. The weights were validated structurally (dims, layout) but never *applied to a field* — and at runtime every exchange was tile-translocated. The canonical pool weights are correct because they were *harvested from a real coupled run* (hence the `per-nproc` pool directories).




Figure 1: **The picture that settles it.** January (month 1) sea-surface temperature as OIFS holds it. Top: movcav4, pool weights — textbook. Bottom: movcav15, offline-engine weights — the ocean is a *tile-translocated jigsaw*: tropical blocks at 55°S, polar blocks on the equator, a straight seam at lon 180 where the mesh node ordering breaks.

Evidence (all on movcav15 itself): OIFS $ci = 0.04$ vs FESOM $a_{ice} = 0.95$ in month 1; 1920 tropical cells $< 5^\circ\text{C}$ and 746 polar cells $> 15^\circ\text{C}$ in OIFS’s January SST (0 and 0 in movcav4); offline application of the exact staged weights to the exact FESOM ice field reproduces OIFS’s received 0.04; global sea-ice death with **no winter regrowth** in both hemispheres; cavity +1.25 K, melt 28 m/yr *regardless of initial state* — the heat arrived through the scrambled exchange every timestep. Every submesh leg since June carried this (masked there by the day-4 crashes); the healthy chunk-1 runs (movcav4/7/8) used pool weights.

Fix (c21fe7c2, option B1 — in the coupling layer, ocp-tool untouched): `permute_feom_to_runtime.py` reorders `feom.*` in grids/masks/areas to runtime order (from the submesh’s own `dist_<nproc>/rpart.out`, verified identical to the `my_list` concatenation FESOM reads); `validate_feom_weights.py` then applies every weight file to the runtime-ordered node latitudes and **aborts the leg** (`couple_fail`) if the median error exceeds 1° (nod2d-ordered weights fail at 42–58°; correct ones pass at $\leq 0.025^\circ$). The validator would have caught the bug the day the engine was introduced. Round-trip tested on compute; the pooled chunk-1 artifacts are fixed in place (backups kept).

1.3 How the two bugs interlocked

Bug I collapsed PISM even in healthy runs $\Rightarrow$ monstrous per-leg geometry increments. Bug II poisoned the runs used to diagnose Bug I: measured on jigsaw-coupled movcav15, “FESOM’s own melt” appeared to be 28 m/yr, which falsely disqualified the DIRECT fix and sent the investigation into the cavity-water/SPP dead end. **Measurements from a corrupted run are worse than no measurements.**

2 Blocker I — OIFS step-0 coastline crash (RESOLVED)

A floating-point NaN struck exactly the coastline cells whose land–sea type flipped. `ocp-tool` half-converted a flipped cell (only mask + soil type), OIFS re-ingested the inconsistency from the re-generated ICMGG and NaN’d in moist physics. Fixed by a masked nearest-neighbour rebuild of each flipped cell’s full column, plus the ICMGG re-ingest (PUTALLFLDS_M in `reresf_part2.F90`). A recurring `voskin` floating-overflow variant shows up state-dependently and was traced to the native-pipeline ICMGG coastal lsm.

3 Blocker II — FESOM cavity-margin blowup

3.1 Two root causes we got wrong

FALSIFIED #1 — “a geostrophically-unbalanced remapped state”. The mesh change plants fresh cavity fronts against salty open water but zeroes the balancing current; initialise the changed-cell velocity in thermal-wind balance instead of zero. That fix was implemented (per-level $\mathbf{u}_g = (1/f\rho_0) \hat{z} \times \nabla p$, using FESOM’s own cavity pressure convention) and **failed its own offline gate**: the discrete ∇p at grid-scale fronts is *noisier* than the nearest-neighbour fill. It ships opt-in only (REMAP_GEOSTROPHIC=1), off by default.

FALSIFIED #2 — “the sub-ice free-surface BC violation”. FESOM keeps $\eta \equiv 0$ under ice, yet the remap gave the 987 newly-sub-ice nodes their old open-ocean ssh (≈ -1.6 m); fixing that will cure the blowup. A **real bug**, and fixed (esm_tools b17fb5c8) — but **not the cure**: the fixed run still died on day 4.4.

3.2 Every hypothesis, and the experiment that killed it

Hypothesis	Experiment	Verdict
Atmosphere / OASIS / coupling	ocean-only standalone FESOM reproduces the crash exactly (day 4, CFLz 9.24)	ruled out
CORE2 forcing shock (standalone artefact)	full-mesh control, <i>same</i> forcing + native restart: stable 21 d, worst CFLz 2.61	ruled out
Remap corrupting the state	remapped vs. source over the 204970 unchanged nodes: $\max \Delta = \mathbf{0.000}$ for temp, salt, ssh, hnode	ruled out (bit-exact)
The restart <i>fill values</i> at changed nodes	five independent fill strategies (donor, coherent-column, shelf-base hold, NN velocity, full balance hygiene)	ruled out (all die $\approx$ day 4)
Pathological thin water columns	submesh min wet column 30 m vs. mother’s 25 m; zero 1–2-layer columns	ruled out (sub-mesh is <i>cleaner</i>)

Hypothesis	Experiment	Verdict
Basal melt flux shock	cavity_flux_ramp_days implemented; melt effectively OFF still dies (mstep 328)	ruled out
Timestep / unresolved w	1200s: CFLz kill day 4.0. 60s: 0 CFLz warnings, dies day 4.4 via η_n . 120s: dies day 8.9	ruled out
Sub-ice free-surface BC	real bug, fixed; run still dies day 4.4	not the cure
Cavity geometry itself	cold start on the same PISM sub-mesh @1200s: STABLE, day 13+, worst CFLz 3.44	ruled out

3.3 The decisive control: a cold start on PISM’s cavity is stable

Same submesh, same forcing, same production timestep — only the initial state differs:

Run	Δt	Initial state	Outcome
full-mesh control	1200s	native spun-up CORE3 restart	stable 21 d, worst CFLz 2.61
PISM cavity, cold	1200s	cold (climatology, at rest)	stable 13+ d, worst CFLz 3.44
PISM cavity, warm	1200s	our remapped restart	dies day 4.0 , CFLz 9.24
PISM cavity, warm	120s	our remapped restart	dies day 8.9 (η_n)
PISM cavity, warm	60s	our remapped restart	dies day 4.4 (η_n , 0 CFLz)

The ocean has no difficulty living in PISM’s cavity at the production timestep. Whatever is wrong is in *our restart*, not in the mesh, the physics, or the coupling.

3.4 What is actually wrong: the dynamics seam

The remap produces an internally inconsistent state: on the **unchanged nodes (98%)** it reproduces the evolved restart **bit-exactly** (a fully developed, balanced dynamic state: $\mathbf{u}$, η , AB history); on the **changed nodes** it substitutes *patched* values. The two regions are mutually incompatible, and the **seam between them** is the instability. Zeroing the dynamics globally removes the crash (the ocean re-derives its own balanced flow from rest) — which is exactly why all five fill strategies failed: each varied the *patch* while leaving the *seam*. Zero-dynamics is not a production fix: discarding the ocean’s circulation every leg is unacceptable.

4 The increment test (2026-07-13) — **TAINTED, to be re-run**


TAINTED: everything in this section was measured on movcav12, whose chunk-1 climate was jigsaw-coupled (Bug II) and whose PISM leg was additionally driven by Bug I forcing. The qualitative machinery findings (the remap carried the increment; the failure mode was 17 fully-opened columns) may survive, but no number here should be trusted until the test is repeated under movcav16.

Cold-starting the ocean on PISM’s cavity sidesteps the one-off CORE3(observed)→PISM swap ($\max|\Delta\text{draft}| = 2711$ m) entirely. That was done: chunk 1 ran a clean model year, PISM ran its 10 years, and the remap was finally asked to carry *a year of spun-up ocean dynamics across a real PISM increment*. It crashed on day 15.7.

	one-off swap (CORE3→PISM)	real increment
CFLz over time	4.46 → 6.43 → 7.25 → 9.24	5.26 → 5.78 → 5.40 → 4.15
	monotonic runaway	peaks, then decays
died at	day ~4	day 15.7
η_n NaN	yes	none — all 209 172 nodes finite
$ \eta > 3$ nodes	3611, basin-wide	17 , one cluster, all changed nodes

All 17 ringing nodes share one signature: $\text{ulev } 17 \rightarrow 1$ — **16 ocean levels opening at once**, i.e. the ice shelf above them removed *entirely*. Their η alternates in sign (+24.4, −17.9, +15.9, +8.3): a grid-scale **checkerboard**, not a basin mode.

4.1 The finding: PISM is collapsing



The cutaway is the clearest statement of the problem: the ice base goes from $-200 \dots -280$ m to **zero (open ocean)** across the whole span -82.2° to -78.2° latitude. **The Ross ice-shelf front retreats ≈ 440 km in ONE 10-yr leg.** The plan view shows this is not a Ross anomaly: the 1399 fully-opened nodes (magenta) ring the *entire* Antarctic margin.

changed nodes this leg	3248	(1.55% of mesh)
$ \Delta_{\text{ulev}} \geq 10$		1582
nodes with the ice removed ENTIRELY		1399
mean $ \Delta_{\text{draft}} $	2.9 m	← meaningless
max $ \Delta_{\text{draft}} $		923 m
PISM floating cells	24 815 $\rightarrow$ 20 093	(−24% in 10 yr)

FALSIFIED #3 — “the per-leg increment is gentle”. ~~Production never has to survive the one-off swap; the per-leg PISM increment is two orders of magnitude smaller (mean $|\Delta_{\text{draft}}| \approx 22$ m) and therefore absorbable.~~ **Wrong, and wrong in a way that mattered** — it motivated the entire “sidestep the swap and the increments will be fine” strategy. The mean is diluted by the 98% of the domain that does not change; the distribution is wildly skewed. PISM is *not* in equilibrium.

4.2 This is not good news

It is tempting to read the numbers optimistically: only 17 of 1399 fully-opened nodes rang; the ocean absorbed $\sim 99\%$; the CFLz transient peaked and *decayed*. All true, and it does mean the remap machinery is performing better than a “crashed at day 15” headline suggests.

But the honest reading is negative. The run still crashed. What we have learned is that **our ocean reacts violently to ice-shelf and cavity changes mid-run**: a handful of columns where the ice vanishes outright is enough to kill a leg. A coupled setup in which the ice sheet is *allowed* to change its cavity every 10 years cannot be one leg’s worth of extreme nodes away from failure. **Absorbing 99% of the change is not a passing grade when the remaining 1% is fatal.**

5 Falsified: the SPP claim

FALSIFIED #4 — “the salt plume parameterisation was off”. ~~spp=false means no brine plume, no HSSW, warm shelf — turn spp on.~~ Checked against the numbers instead of the figure: in Xiaojie’s CORE3 runs the **spp-OFF cavity is the coldest** (thermal driving +0.177 K vs +0.246 with spp on), and CORE3 standalone keeps a cold cavity *with our exact setting*. The warm cavity in movcav12/15 was Bug II (the jigsaw), not spp. The claim was made from reading a shelf-transect figure without checking whether the signal reached the cavity. Config change reverted.

6 Where to go from here

movcav16 is the clean end-to-end test: fixed weights + warm-tracer init + DIRECT melt coupling. Verification checklist for chunk-1 (each item targets one of this week’s failure modes):

1. OIFS January SST map is physical (the jigsaw figure *is* the test);
2. OIFS $c_i \approx$ FESOM `a_ice` (was 0.04 vs 0.95);
3. sea ice survives the winter in both hemispheres (was: dead by April, no regrowth);
4. FESOM cavity melt ≈ 2 m/yr (was 28); cavity thermal driving ≈ 0.2 K (was 1.25);
5. PISM leg: floating-area change per decade $\ll 24\%$;
6. then chunk-3: the *real* increment test, finally on clean evidence.

6.1 movcav16 launch log (2026-07-14) — and Bug III

Four attempts. Attempts 2 and 3 exposed a **third ordering bug**, sibling of Bug II.

- **Attempt 1** (job 26245041): healthy to day 120 (worst CFL_z 2.41, EVP active, ~ 17 d/min) – killed by a launcher bug of ours (`env_fesom.py` looked up `oce2ice_method` under the setup name instead of `config['fesom']`; fixed in `esm_tools` d5a8039d). The teardown of the crashed launcher cancelled the healthy compute job.
- **Attempt 2** (job 26245420): OIFS floating overflow in the convection scheme (`cubasen.F90:453`, rank 29) at hour 458 (19 January); the watchdog killed the job correctly. FESOM healthy; work-dir `rmf_*feom*` bit-identical to the corrected pool weights, so Bug II's fix *was* in effect. `BETAMFBC=-1` already set (not the `qmfixer` overflow). Forensics: `movcav16_crash1_forensics/`.
- **Attempt 3** (job 26246140): same routine, `cubasen.F90:453/454`, but at *hour 19–20* and on several ranks at once (overflow + divide-by-zero, ranks 10/27). Same crash site, times spread over hour 19 / hour 458 / never (attempt 1's 120 days) $\Rightarrow$ a *startup shock* that nondeterministically detonates marginal convection columns.
- **Bug III — the OASIS coupling restart was compound-scrambled.** The chunk-1 pool `rstos.nc` (OIFS's SST/ice for the first coupling interval, before FESOM's first real field arrives) was built by `remap_oasis_restart.py`, whose `map_nod` gather lives in `nod2d` space — but the harvested input is *runtime*-ordered on the old mesh, and the runtime-ordered weights need the output *runtime*-ordered on the new mesh. Both permutes were missing: the emitted field matched *neither* ordering. Measured: pool `sst_feom` correlated +0.15 with the restart's own level-1 temperature (and -0.11 with $|\text{lat}|$); the corrected transform (un-permute via old `rpart.out` $\rightarrow$ gather $\rightarrow$ re-permute via new `rpart.out`) correlates +0.995 (and -0.95 with $|\text{lat}|$). So OIFS's first hour saw 305 K water under polar air — a one-interval global jigsaw shock. Fix: `esm_tools` abe56ee2 (script + caller pass both `rpart.out`); pool `rstos.nc` rebuilt from the harvested full-mesh original (backup: `rstos_scrambled_bak.nc`); `rstas.nc` is A096-only, unaffected. *This closes the “transient-only, deferred” loose end — the transient was not harmless.*
- **Attempt 4** (job 26246292): corrected `rstos.nc` verified in the work dir. Survived to *day 44* (longest of the crashing attempts — consistent with Bug III's fix helping) then died in the *surface scheme*: `vsurf_mod.F90:280`, `FOEEW(ZTSK)` overflow = a runaway skin temperature, on a tile the OIFS source itself flags (“Tsk ... overshoots to very large values ... for non-physical tile fractions”). FESOM healthy throughout ($CFL_z \leq 1.8$).
- **The remaining suspect: the native-pipeline ICMGG.** `fort.4` is byte-identical between surviving `movcav15` and crashing `movcav16`; WAM tables, `plit`, masks identical too; all mask cross-checks (`lsm` $\leftrightarrow$ `plit` $\leftrightarrow$ OASIS mask $\leftrightarrow$ weight coverage) are clean; the corrected weights deliver OIFS geographically-correct fields (max 0.24°). What `movcav16` does *first* is combine the native ICMGG (74 hardened Antarctic coastal cells) with *physically correct* polar coupled fields — under `movcav15`'s jigsaw, those cells never saw real polar values (`ci` ≈ 0.04 , ice tile effectively off). Every OIFS `libsurf/convection` crash ever observed (`voskin/vsurf/cubasen`) is on a native-ICMGG run; `movcav4/7` with the clean CORE3 file never crashed; and the chunk-3 swap test already *proved* the ICMGG mechanism there.
- **Attempt 5: an invalid test with a lesson.** `ICMGG_INIT` $\rightarrow$ clean `ICMGGab45INIT_CORE3`, everything else unchanged — died at *step 0* in `radaca.F90:810` (radiation), dividing by a *zero half-level temperature*, on a far-south rank. Cause (verified by cell count): the pool *plit* still flips land-ice on **62 cells the clean `lsm` calls ocean** — the degenerate-skin-T trigger the `plit` machinery exists to avoid. **ICMGG and `plit` must be swapped as a pair.** (With the native `lsm` the `plit` conflict count is 0.)

- **Attempt 6: same step-0 death $\Rightarrow$ the native surface set is REQUIRED.** Clean ICMGG *plus* `ECE_LANDICE=.false.` and no `plit` (movcav4’s exact atmosphere) died *identically* to attempt 5: `radaca.F90:810`, step 0, far-south rank, zero half-level temperature. So the zero-T column is not the `plit` flip — it is the **62 difference-zone cells themselves**: ocean per the clean lsm, but never delivered by the `feom`→A096 mapping, so their coupled SST stays **0 K** (the original “degenerate-skin-T trigger”). Conclusion: on the PISM submesh there is *no valid static-lsm configuration*; the native, coastline-following surface set (`lsm+slt+plit+A096 mask+weights`) is the only one that can even initialize — which is also what the moving-cavity mission requires. The swap test cannot run, and the ICMGG is *not* convicted for the mid-run crashes.
- **Attempt 7 (running): native set + OASIS EXPOUT instrumentation.** Every o2a exchange (`sst/sie/snt/ist/sia/u/v`) is dumped to netcdf (`oasis_expout_o2a` runscript flag, `esm_tools 8b5dfca2`). On the next crash the last records identify the poisoned cell and field directly. Suspicion is guided by the FP-error family: FOEEW-style saturation math blows up on near-zero or near-32 K temperatures, so a near-zero `A_SST/A_Ice_temp` at some cell, some hour, is the expected signature. Also cleared meanwhile: raw weight-sums on all active A096 cells are exactly 1.0 (no under-delivered tiny-sum cells). Crash-rank geography: attempts 2–4 died far-north (ranks 10/27/29/50 of 384), attempts 5/6 far-south (367/381) — polar always, tropics never.
- Launch hygiene from attempts 2–3: launch from the iterative master (`spinup_couplePI_core3.yaml`; the `modell` sub-runscript alone dies with `KeyError: modell`), and `use_venv: False` so non-tty launches cannot hang on the `venv` prompt.

If attempt 5 convicts the ICMGG, the durable fix is structural: the atmosphere-side surface description (ICMGG `lsm/slt` + A096 OASIS mask + `feom` weights + WAM tables) must change *as a consistent set* or not at all — today the WAM tables are still the old coastline (byte-identical to movcav4’s) even where the lsm moved.

7 Bugs found and fixed along the way

7.1 In the remap

Unbounded extrapolation below the old seabed (`esm_tools eef4f6db`). The first attempt at the increment leg died at **mstep 1**: the remapped restart contained -45.02°C , a value present *nowhere* in the source (wet range $[-2.32, 41.72]$, dry cells 0.0). The remap *manufactured* it. The column shows the mechanism: 1.22, -0.99 , -3.19 , -5.39 , -8.69 , -13.09 , -18.6 , -25.2 , -34.01 , **-45.02** — smooth and accelerating. For columns that get *deeper*, the remap fitted a straight line through the last two old levels and continued it downward with no clamp. Fixed by taking such water from the nearest old node actually wet at that depth. **Never hit before** because every previous remap went mother $\rightarrow$ submesh (a strict *subset*); ice *retreating* grows the mesh and exercises this path for the first time.

Also: sub-ice `ssh/hbar` zeroing (`b17fb5c8`); a **size guard** so the remap refuses a restart that does not match the old mesh, instead of silently writing garbage (`3b44a704`).

7.2 In the coupling plumbing

Making chunk 1 a *submesh* leg pushed a submesh leg through `couple_out` for the first time. Every previous run either ran chunk 1 on the *full mesh* (where these paths are trivially correct) or died early in chunk 3 (before reaching the end of a leg). That exposed:

- **OASIS GSSPOS abort at ~day 300.** Its conservation correction is built from the summed source/target ratio; when the summed target residual is ≈ 0 while the source residual is small-but-nonzero, it explodes. Only the heat-to-ice flux (`A_Q_ice` $\rightarrow$ `heat_ico`) used it. Fixed: new `gauswgt_gsmart` transform (3e94d4da).
- **The `namcouple` feom dimension had NEVER been patched, on any leg.** `fix_namcouple_feom_dim` read `#{MAXMESH_DIR_fesom}`, which is assigned *inside* `build_submesh` — a different subjob’s shell. Always unset $\Rightarrow$ guard refused $\Rightarrow$ silent no-op. Fixed (3cf29e9e).
- **`fesom2ice` loaded the static full mesh** while FESOM’s output is on the submesh (broadcast error) (63cfcc91).
- **`previous_submesh` is read but never created** — anywhere. Every leg from chunk 5 on would have remapped against the wrong mesh (3b44a704).
- **A failed `couple_in` did not stop the workflow.** The coupling function is backgrounded, the observe job is not its parent (so it can never see the exit code), and `observe.py` returns an *empty* error list for every non-compute job. The next leg started on an un-remapped restart. Fixed with a `.coupling_failed` sentinel (3b44a704).

The pattern: the guards were right, and the plumbing around them defeated them. `remap_oasis_restart` correctly refused a mis-sized restart — and the run proceeded anyway.

8 Assets

- **A 2-minute offline reproducer:** standalone ocean-only FESOM on the PISM-cavity sub-mesh, 14 nodes, crashes in ~ 2 min (vs. the 45-node coupled system). Requires the `FESOM_COUPLED=OFF` build at the current SHA.
- **A stable control:** full mesh + native restart + same forcing (21 d, CFLz 2.61).
- **Chunk-1 artifacts pooled** at `input/fesom2/pism_cavity_ini/` (540 MB), so chunk 1 needs no `couple_in` and `esm_runscripts` can be run straight from the login node.
- `cavity_flux_ramp_days` (FESOM `&run_config`, default 0 = off, bit-identical).
- `verify_balance.py`: offline balance/stability gate ($\min N^2$, $\max |\nabla \cdot (H\bar{\mathbf{u}})|$, CFL, Ri) — it caught the geostrophic ∇p init being worse than the NN baseline *before* it cost a model run.

9 Method note

Three root causes were asserted with more confidence than the evidence supported, and all three were later falsified by a control that was cheap and should have come *first*:

1. “geostrophic imbalance” — killed by the **cold start** (is it the geometry or our state?);
2. “the sub-ice ssh BC” — killed by simply **running the fix**;
3. “the increment is gentle (≈ 22 m)” — killed by **plotting the distribution** instead of trusting a mean over a domain that is 98% unchanged.

The third is the sharpest: the 22 m figure was computed correctly and reported honestly, and it was still misleading enough to steer the strategy for a day. **Two figures did in ten minutes what a summary statistic obscured for a day.**